

RBF Finite-Difference Discretization of Anisotropic Diffusion: Theory and Implementation

Dennis Corraliza

6/21/2026

Abstract

This note derives the Radial Basis Function Finite Difference (RBF-FD) method for the anisotropic diffusion equation $-\nabla \cdot (A \nabla u) = f$ on a bounded domain $\Omega \subset \mathbb{R}^2$, with mixed Dirichlet/Neumann boundary conditions, and shows precisely how each step of the derivation corresponds to a function in the accompanying Python implementation (modules `assembly.py`, `domain.py`, `rbf.py`, `boundary.py`, `stencils.py`, and `geometry.py`).

Contents

1	Problem statement	2
2	Radial basis function interpolation	2
2.1	Local RBF interpolant	2
3	RBF-FD differentiation weights	3
3.1	The linear functional viewpoint	3
3.2	Least-squares variant	4
4	The anisotropic diffusion operator applied to each kernel	4
5	Global assembly	5
5.1	Interior rows	5
5.2	Dirichlet rows	5
5.3	Neumann rows	6
5.4	Right-hand side	6
6	The pure-Neumann null space and anchoring	6
7	Solving and evaluating the discrete solution	6
8	Convergence diagnostics	7
9	Manufactured solutions	7
10	Node generation strategies	8

1 Problem statement

We seek $u : \Omega \rightarrow \mathbb{R}$ satisfying the anisotropic diffusion equation

$$-\nabla \cdot (A \nabla u) = f \quad \text{in } \Omega, \quad (1)$$

where $A \in \mathbb{R}^{2 \times 2}$ is a constant, symmetric positive-definite diffusion tensor and f is a prescribed source term. The boundary $\partial\Omega$ is partitioned into Dirichlet and Neumann pieces:

$$u = g_D \quad \text{on } \Gamma_D \subset \partial\Omega, \quad (2)$$

$$\frac{\partial u}{\partial n} \equiv n \cdot \nabla u = g_N \quad \text{on } \Gamma_N \subset \partial\Omega, \quad (3)$$

where n is the outward unit normal. When $A = I$, (1) reduces to the standard Poisson equation $-\Delta u = f$.

In the code, the geometric domain is either a square $[0, L]^2$ or a disk of radius R (`geometry.py`), the tensor A is supplied directly by the user, and the boundary type on each side/arc is encoded by the list `btype` (e.g. `['dirichlet', 'neumann', 'dirichlet', 'neumann']` for a square, ordered counterclockwise starting at $y = 0$).

2 Radial basis function interpolation

2.1 Local RBF interpolant

Let $\{\mathbf{x}_j\}_{j=1}^N$ be the full node set discretizing Ω . For a fixed node $\mathbf{x}_i$, let $S_i = \{j_1, \dots, j_n\} \subset \{1, \dots, N\}$ index its *stencil*: the n nearest neighbors of $\mathbf{x}_i$ (including $\mathbf{x}_i$ itself), determined by a k -nearest-neighbor search. We build a local interpolant of a function u over the stencil using a radial kernel φ :

$$u(\mathbf{x}) \approx \sum_{j \in S_i} \lambda_j \varphi(\|\mathbf{x} - \mathbf{x}_j\|) + \sum_{\ell=1}^m \mu_\ell p_\ell(\mathbf{x}), \quad (4)$$

where $\{p_\ell\}_{\ell=1}^m$ is an (optional) polynomial basis used to *augment* the kernel interpolant, with coefficients constrained by

$$\sum_{j \in S_i} \lambda_j p_\ell(\mathbf{x}_j) = 0, \quad \ell = 1, \dots, m. \quad (5)$$

Two kernels are used in the implementation (`rbf.py`):

$$\varphi_{\text{cubic}}(\mathbf{p}) = \|\mathbf{p}\|^3, \quad \varphi_{\text{gauss}}(\mathbf{p}) = \exp(-(\varepsilon \|\mathbf{p}\|)^2), \quad (6)$$

where $\mathbf{p} = \mathbf{x} - \mathbf{x}_j$ is a displacement vector and ε is a user-chosen shape parameter for the Gaussian. The polynomial basis used for augmentation is linear,

$$\{p_1, p_2, p_3\}(x, y) = \{1, x, y\}, \quad (7)$$

so that the interpolant exactly reproduces constants and linear functions whenever augmentation is enabled.

Code correspondence.

- `rbf.phi_cubic(p), rbf.phi_gauss(p, eps) ↔ (6).`
- `rbf.poly_basis(p) ↔ (7).`
- `stencils.knn_list(P, num_stencil_nodes)` builds S_i for every node.
- `domain.PDEDomainContext.augmentation` is the boolean flag toggling the second sum in (4).

3 RBF-FD differentiation weights

3.1 The linear functional viewpoint

RBF-FD does not differentiate the interpolant (4) symbolically; instead, it seeks weights w_j such that a linear operator $\mathcal{L}$ applied at $\mathbf{x}_i$ is approximated directly as a weighted sum of *nodal values*:

$$(\mathcal{L}u)(\mathbf{x}_i) \approx \sum_{j \in S_i} w_j u(\mathbf{x}_j). \quad (8)$$

The weights are chosen so that (8) is *exact* whenever u is any kernel function centered at a stencil node, and (if augmentation is used) exact for every polynomial in the augmentation basis. This yields the linear system

$$\underbrace{\begin{bmatrix} M & P \\ P^\top & 0 \end{bmatrix}}_{\text{generalized Gram matrix}} \begin{bmatrix} \mathbf{w} \\ \boldsymbol{\mu} \end{bmatrix} = \begin{bmatrix} \mathbf{b} \\ \mathbf{0} \end{bmatrix}, \quad (9)$$

where

$$M_{jk} = \varphi(\mathbf{x}_j - \mathbf{x}_k), \quad j, k \in S_i, \quad (10)$$

$$P_{j\ell} = p_\ell(\mathbf{x}_j), \quad j \in S_i, \ell = 1, \dots, m, \quad (11)$$

$$b_j = (\mathcal{L}\varphi)(\mathbf{x}_i - \mathbf{x}_j), \quad j \in S_i. \quad (12)$$

The vector $\mathbf{b}$ holds the operator $\mathcal{L}$ applied *exactly* (in closed form) to the kernel, evaluated at the displacement from the stencil center to each stencil node, rather than any discrete approximation. Solving (9) for $\mathbf{w}$ gives the desired differentiation weights at node $\mathbf{x}_i$.

Code correspondence. This is exactly what `assembly.local_weights_solve(context,i)` and `assembly.local_grad_solve(context,i)` implement, for $\mathcal{L} = \nabla \cdot (A\nabla \cdot)$ and $\mathcal{L} = \nabla$ respectively:

- The matrix `M` in the code is exactly M in (12), built from `context.phi`.
- The vector `b` (or matrix `b_grad`, with one column per spatial dimension) is built from `context.laplacian_phi` (resp. `context.grad_phi`), which are the closed-form expressions $\nabla \cdot (A\nabla\varphi)$ and $\nabla\varphi$ derived in Section 4 below.
- The block matrix `Pmat` and the block-augmented system built via `np.block` are exactly the augmented system in (9) when `context.augmentation` is `True`.
- `np.linalg.solve(M, b)` solves (9) and the function returns only the first `num_nodes` entries of the solution, i.e. $\mathbf{w}$ (discarding the polynomial Lagrange multipliers $\boldsymbol{\mu}$).

3.2 Least-squares variant

An alternative formulation avoids solving the square system (9) directly and instead seeks weights that best reproduce $\mathcal{L}$ in a least-squares sense over a richer set of *auxiliary evaluation centers* $\{\mathbf{c}_k\}_{k=1}^K$ scattered around $\mathbf{x}_i$ (typically more centers than stencil nodes, $K > n$):

$$\min_{\mathbf{w}} \left\| \underbrace{\left[\varphi(\mathbf{c}_k - \mathbf{x}_j) \right]_{k,j}}_{=: M \in \mathbb{R}^{K \times n}} \mathbf{w} - \underbrace{\left[(\mathcal{L}\varphi)(\mathbf{x}_i - \mathbf{c}_k) \right]_k}_{=: \mathbf{b} \in \mathbb{R}^K} \right\|_2^2, \quad (13)$$

subject to the same polynomial reproduction constraints (5) enforced exactly (as equality rows appended to the system, rather than penalized in the least-squares objective). The motivation for this variant is robustness: anchoring the fit to a denser, more uniformly distributed cloud of auxiliary points reduces sensitivity to clustering or anisotropy in the stencil node positions themselves.

The auxiliary centers $\{\mathbf{c}_k\}$ are generated as a near-uniform polar grid of *num_rings* concentric rings (spacing $h = r/\text{num}_r\text{ings}$, with r the stencil radius) plus a center point, with the number of angular samples per ring chosen so that the arc-length spacing on each ring is approximately h as well:

$$\mathbf{c}_k = (\rho_i \cos \theta_{ij}, \rho_i \sin \theta_{ij}), \quad \rho_i = ih, \quad \theta_{ij} = \frac{2\pi j}{n_i}, \quad n_i = \max\left(1, \text{round}\left(\frac{2\pi\rho_i}{h}\right)\right). \quad (14)$$

Code correspondence.

- `rbf.generate_grid_2d(r, k)` implements (14), using `rbf.rad_to_euc` to convert each (ρ_i, θ_{ij}) to Cartesian coordinates.
- `assembly.local_weights_ls(context, i)` and `assembly.local_grad_ls(context, i)` build M and $\mathbf{b}$ exactly as in (13), append the polynomial constraint rows, and call `np.linalg.lstsq` to solve.
- Dispatch between the direct system (9) and the least-squares system (13) is controlled by `context.center_rings`: `None` selects the direct solve, any integer value selects the least-squares variant with that many rings *num_rings*.

4 The anisotropic diffusion operator applied to each kernel

The entries b_j in (12) require a closed-form expression for $\nabla \cdot (A \nabla \varphi)$, evaluated at the displacement $\mathbf{p} = \mathbf{x}_i - \mathbf{x}_j$. Writing $A = \begin{pmatrix} A_{11} & A_{12} \\ A_{12} & A_{22} \end{pmatrix}$ and $r = \|\mathbf{p}\|$, $\mathbf{p} = (x, y)$:

Cubic kernel. For $\varphi(\mathbf{p}) = r^3$,

$$\nabla \cdot (A \nabla \varphi)(\mathbf{p}) = 3r \left(\text{tr}(A) + \frac{\mathbf{p}^\top A \mathbf{p}}{r^2} \right). \quad (15)$$

This expression is regular as $r \rightarrow 0$ in the sense that the true value of the operator at $\mathbf{p} = 0$ is 0; the implementation returns 0 directly whenever r falls below a tolerance, avoiding the removable $1/r^2$ singularity in (15) without altering the limiting value.

Gaussian kernel. For $\varphi(\mathbf{p}) = \exp(-(\varepsilon r)^2)$,

$$\nabla \cdot (A \nabla \varphi)(\mathbf{p}) = 2\varepsilon^2 \varphi(\mathbf{p}) \left(2\varepsilon^2 \mathbf{p}^\top A \mathbf{p} - \text{tr}(A) \right). \quad (16)$$

No singularity occurs here since φ and its derivatives are smooth everywhere.

Polynomial augmentation basis. Since $\nabla \cdot (A \nabla p_\ell) \equiv 0$ for $p_\ell \in \{1, x, y\}$ (all second derivatives vanish), the corresponding entries in the augmented right-hand side of (9) are identically zero, independent of A .

Gradients. The two kernel gradients used to build Neumann boundary rows (Section 5.3) are

$$\nabla \varphi_{\text{cubic}}(\mathbf{p}) = 3r \mathbf{p}, \quad \nabla \varphi_{\text{gauss}}(\mathbf{p}) = -2\varepsilon^2 \varphi_{\text{gauss}}(\mathbf{p}) \mathbf{p}. \quad (17)$$

Code correspondence.

- `rbf.anisotropic_diffusion_phi_cubic(p, A, tol) ↔ (15);`
`rbf.anisotropic_diffusion_phi_gauss(p, A, eps) ↔ (16).`
- `rbf.anisotropic_diffusion_poly(p, A, tol)` returns 0.0, matching the polynomial-basis observation above.
- `rbf.grad_phi_cubic(p), rbf.grad_phi_gauss(p, eps) ↔ (17).`
- `assembly.set_rbf_func` wires the chosen kernel, gradient, and anisotropic-diffusion closures (with A fixed) onto `context.phi`, `context.grad_phi`, and `context.laplacian_phi` respectively, dispatching on the string `basis` ('cubic' or 'gaussian').

5 Global assembly

5.1 Interior rows

For every node $\mathbf{x}_i$ in the domain, the local weight vector $\mathbf{w}^{(i)} \in \mathbb{R}^{|S_i|}$ computed by solving (9) (or (13)) is scattered into row i of a dense global matrix $W \in \mathbb{R}^{N \times N}$:

$$W_{ij} = \begin{cases} w_j^{(i)} & j \in S_i, \\ 0 & j \notin S_i. \end{cases} \quad (18)$$

At this stage, W approximates $\nabla \cdot (A \nabla \cdot)$ at every node, ignoring boundary conditions entirely.

Code correspondence. `assembly.global_weights(context, in_boundary, normal_vec)` loops over all stencils, dispatches to `local_weights_solve` or `local_weights_ls` depending on `context.center_rings`, and assembles (18).

5.2 Dirichlet rows

For a node $\mathbf{x}_i \in \Gamma_D$, row i of W is overwritten with the identity row,

$$W_{ij} = \delta_{ij}, \quad i \in \Gamma_D, \quad (19)$$

so that the linear system enforces $u(\mathbf{x}_i) = g_D(\mathbf{x}_i)$ exactly (via the matching right-hand-side entry, Section 5.4).

5.3 Neumann rows

For a node $\mathbf{x}_i \in \Gamma_N$ with outward unit normal n_i , row i is overwritten with directional-derivative weights, built from the *local gradient* weight matrix $W_{\text{grad}}^{(i)} \in \mathbb{R}^{|S_i| \times 2}$ (the analog of (9) for $\mathcal{L} = \nabla$):

$$W_{ij} = \left(W_{\text{grad}}^{(i)} n_i \right)_j, \quad j \in S_i, \quad i \in \Gamma_N, \quad (20)$$

so that row i of the assembled system enforces $n_i \cdot \nabla u(\mathbf{x}_i) \approx g_N(\mathbf{x}_i)$, matching (3).

Code correspondence. `assembly.boundary_to_weights` implements (19) and (20) in place, dispatching per node on the label returned by `in_boundary(p)` (itself built in `assembly.set_boundary_func` from `boundary.in_square_boundary` / `boundary.in_circular_boundary`, with normals from `boundary.square_normal` / `boundary.disk_normal`).

5.4 Right-hand side

The right-hand-side vector $\mathbf{F} \in \mathbb{R}^N$ is assembled nodewise,

$$F_i = \begin{cases} f(\mathbf{x}_i) & \mathbf{x}_i \text{ interior,} \\ g_D(\mathbf{x}_i) & \mathbf{x}_i \in \Gamma_D, \\ g_N(\mathbf{x}_i) & \mathbf{x}_i \in \Gamma_N, \end{cases} \quad (21)$$

so that the assembled linear system $W\mathbf{u} = \mathbf{F}$ is exactly the discrete analog of (1)–(3).

Code correspondence. `assembly.right_hand_side(context, f, g, in_boundary)` implements (21), where the single callable `g` passed in already dispatches internally to the correct side/arc function (see `boundary.square_boundary`, `boundary.circular_boundary`).

6 The pure-Neumann null space and anchoring

If $\Gamma_D = \emptyset$ (every boundary node is Neumann), the continuous problem (1)–(3) determines u only up to an additive constant, since $\nabla \cdot (A\nabla c) = 0$ and $n \cdot \nabla c = 0$ for any constant c . The discrete matrix W inherits this null space and is singular. Three regularization strategies are implemented:

$$\textbf{mean: } \sum_i u_i = 0, \quad (\text{replace the last row of } W \text{ by } \mathbf{1}^\top, \text{ set } F_{\text{last}} = 0), \quad (22)$$

$$\textbf{pin: } u_0 = 0, \quad (\text{replace the first row of } W \text{ by } \mathbf{e}_0^\top, \text{ set } F_0 = 0), \quad (23)$$

$$\textbf{project: } W \leftarrow \Pi W \Pi, \quad \mathbf{F} \leftarrow \Pi \mathbf{F}, \quad \Pi = I - \frac{1}{N} \mathbf{1} \mathbf{1}^\top. \quad (24)$$

Code correspondence. `assembly.is_pure_neumann(context, in_boundary)` detects this case; `assembly.anchor_system(W, f, method)` implements the three strategies in (24), selected by the string `method` ('mean', 'pin', or 'project').

7 Solving and evaluating the discrete solution

The fully assembled (and, if necessary, anchored) linear system

$$W\mathbf{u} = \mathbf{F} \quad (25)$$

is solved directly for the nodal solution vector $\mathbf{u} \in \mathbb{R}^N$, approximating $u(\mathbf{x}_i)$ at every node.

Code correspondence. `assembly.rbf_fd_solve(W, F)` calls `np.linalg.solve(W, F)`, implementing (25). The full pipeline — nodes and stencils, kernel and boundary configuration, weight assembly, boundary imposition, right-hand-side assembly, and anchoring — is orchestrated end-to-end by `assembly.rbf_fd_system(...)`, which returns a populated `domain.PDEDomainContext` carrying A , W , and $\mathbf{F}$ for later use by `rbf_fd_solve`.

8 Convergence diagnostics

For manufactured-solution problems (Section 9), the discrete error is reported in both the maximum norm and a discrete L^2 norm:

$$\|e\|_\infty = \max_i |u_i - u_{\text{exact}}(\mathbf{x}_i)|, \quad \|e\|_2 = \frac{1}{\sqrt{N}} \left(\sum_{i=1}^N (u_i - u_{\text{exact}}(\mathbf{x}_i))^2 \right)^{1/2}, \quad (26)$$

together with the condition number $\kappa(W) = \|W\| \|W^{-1}\|$ as a diagnostic for ill-conditioning of the assembled system (which tends to grow as stencils shrink or as *num_rings*, ε , or node spacing are pushed to extremes).

Code correspondence. These are exactly the quantities computed and printed in `report_and_graph(context, u_exact)`: `np.linalg.cond(W)` for $\kappa(W)$, `np.max(error)` for $\|e\|_\infty$, and `np.linalg.norm(error)/np.sqrt(len(P))` for $\|e\|_2$, alongside contour and surface plots of the approximate solution, exact solution, and pointwise error.

9 Manufactured solutions

To validate the method, exact solutions u_{exact} are chosen in closed form, and the corresponding forcing term f and boundary data g_D, g_N are derived analytically so that u_{exact} satisfies (1)–(3) exactly (the *method of manufactured solutions*). For example, the mixed polynomial–trigonometric solution

$$u_{\text{exact}}(x, y) = (L - x) \left(y - \frac{y^2}{L} \right) + c(L - y) \frac{x(L - x)}{L} + \text{Amp} \sin\left(\frac{\alpha\pi x}{L}\right) \sin\left(\frac{\beta\pi y}{L}\right), \quad c = 2, \quad (27)$$

yields, after applying $-\nabla \cdot (A \nabla \cdot)$ to (27) and collecting polynomial, trigonometric, and mixed-derivative contributions, a forcing term f with three additive pieces:

$$f_{\text{poly}} = -2cA_{11} \frac{L - y}{L} + 2A_{12} \left(\frac{2cx}{L} + \frac{2y}{L} - \frac{c + 1}{L} \right) - 2A_{22} \frac{L - x}{L}, \quad (28)$$

$$f_{\text{trig}} = -\text{Amp} \left(A_{11}k_x^2 + A_{22}k_y^2 \right) \sin(k_x x) \sin(k_y y), \quad (29)$$

$$f_{\text{mixed}} = 2 \text{Amp} A_{12} k_x k_y \cos(k_x x) \cos(k_y y), \quad (30)$$

with $k_x = \alpha\pi/L$, $k_y = \beta\pi/L$, and $f = f_{\text{poly}} + f_{\text{trig}} + f_{\text{mixed}}$, together with Dirichlet data read off directly from (27) on each edge.

Code correspondence. `example_3(Amp, modes, A, L)` returns the tuple `(f, g, btype, u_exact)` implementing exactly (27)–(30): `u_exact` evaluates (27), `f` evaluates the sum of (30), and `g` is the list of four one-variable boundary functions obtained by restricting (27) to each edge of the square.

10 Node generation strategies

The implementation supports several node distributions (`geometry.py`):

- **Uniform tensor-product grid** (`uniform_square`): a regular $N_x \times N_y$ Cartesian grid over $[0, L]^2$, boundary and interior nodes undifferentiated.
- **Uniform interior + explicit boundary** (`uniform_int_square`): interior nodes from a regular grid strictly inside $[0, L]^2$, concatenated with N_b equally spaced nodes per edge (corners de-duplicated); returns the interior node count alongside the array so that boundary/interior can be distinguished downstream if needed.
- **Chebyshev–Gauss–Lobatto grid** (`cheby_square`): tensor-product nodes $x_j = \frac{L}{2}(1 + \cos(j\pi/N_x))$, clustering near the edges, as is standard for spectral collocation methods.
- **Circular domain** (`circular_geometry`): interior nodes from a Cartesian grid filtered to the open disk $x^2 + y^2 < R^2$, plus boundary nodes placed exactly on the circle at N_{bound} equally spaced angles.

In every case, the resulting node array $P \in \mathbb{R}^{N \times 2}$ is the sole input to `domain.PDEDomainContext`, after which stencils (`stencils.knn_list`) and RBF-FD weights are computed purely from the node positions, independent of how they were generated.